

The limitations of MB Kalman filtering and DD state estimation motivate a hybrid approach that exploits the best of both worlds; i.e., the soundness and low complexity of the

Recent years have witnessed remarkable successes of deep neural networks (DNNs) in real-life applications. These data-driven (DD) parametric models were shown to be able to catch the subtleties of complex processes and replace the need to explicitly characterize the domain of interest [18], [19]. Therefore, an alternative strategy to implement state estimation—without requiring explicit and accurate knowledge of the SS model—is to learn this task from data using deep learning. DNNs such as recurrent neural networks (RNNs)—i.e., long short-term memory (LSTM) [20] and gated recurrent units (GRUs) [21]—and attention mechanisms [22] have been shown to perform very well for time series related tasks mostly in interactable environments, by training these networks mostly in an end-to-end, model-agnostic manner from a large quantity of data. Nonetheless, DNNs do not incorporate many trainable parameters and lack the interpretability of MB methods, sequences [23] and lack the interpretability of MB methods. These constraints limit the use of highly parameterized DNNs for real-time state estimation in applications embedded in hardware-limited mobile devices such as drones and vehicles.

The common thread among these aforementioned filters is that they are model-based (MB) algorithms; namely, they rely on accurate knowledge and modeling of the underlying dynamics as a fully characterized SS model. As such, the performance of these MB methods critically depends on the validity of the domain knowledge and model assumptions. MB filtering algorithms designed to cope with some level of uncertainty in the SS models, e.g., [15]-[17], are rarely capable of achieving the performance of MB filtering with full domain knowledge, and rely on some knowledge of how much their postulated model deviates from the true one. In many practical use cases the underlying dynamics of the system is non-linear, complex, and difficult to accurately characterize as a tractable SS model, in which case degradation in performance of the MB state estimators is expected.

unscented Kalman filter (UKF) [10]. Methods based on sequential Monte-Carlo (MC) sampling, such as the family of particle filters (PFs) [11]–[13], were introduced for state estimation in non-linear, non-Gaussian SS models. To date, the KF and its non-linear variants are still widely used for tracking filtering in numerous real world applications involving online filtering.

I. INTRODUCTION

The filtering problem is at the core of real-time tracking. Here, one must provide an instantaneous estimate of the state

B. Data-Aided Filtering Problem Formulation

SS models are studied in the context of several different tasks; these tasks are different in their nature, and can be roughly classified into two main categories: observation and approximation. The first category of tasks is called *state recovery*: observations are given as parts of the observed signal y_t . This can correspond, for example, to the prediction of future observations in a given block via input-output relations in a block of hidden state vectors, given a block of observations in a given block via input-output relations; the generation of missing observations given past observations; the recovery of future observations of a hidden state vector x_t . This family of state recovery tasks includes offline recovery, also referred to as smoothing, filtering; i.e., *online* recovery of x_t from past and current noisy observations of y_t . For a given x_0 , filtering involves the design of a mapping from y_t to \hat{x}_t , $\forall t \in \{1, 2, \dots, T\} \equiv T$, where T is the time horizon.

$$(2) \quad \cdot^t \mathbf{x} \cdot \mathbf{H} = (\mathbf{x})^t \mathbf{h} \quad \mathbf{E} \cdot \cdot^t \mathbf{x} = (\mathbf{E})^t (\mathbf{x})$$

In (1a), x_t is the latent state vector of the system at time t , which evolves from the previous state x_{t-1} , by a (possibly) non-linear, state-evolution function $f(\cdot)$ and by an AWGN non-linear, state-evolution function $\mathbf{g}(\cdot)$, and by a (possibly) non-linear, state-observation function $\mathbf{h}(\cdot)$. In (1b), y_t is the vector of observations at time t , which is generated from the current latent state vector by a (possibly) non-linear observation (emission) mapping $\mathbf{h}(\cdot)$ corrupted by AWGN v_t , with covariance \mathbf{R} . For the special case where the evolution or the observation transformations are linear, there exist matrices \mathbf{F} , \mathbf{H} such that

$$\mathbf{x}_t = \mathbf{f}(\mathbf{x}_{t-1}) + \mathbf{w}_t, \quad \mathbf{w}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}), \quad \mathbf{x}_t \in \mathbb{R}^m, \quad (1a)$$

$$\mathbf{y}_t = \mathbf{h}(\mathbf{x}_t) + \mathbf{v}_t, \quad \mathbf{v}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{R}), \quad \mathbf{y}_t \in \mathbb{R}^n. \quad (1b)$$

We consider dynamical systems characterized by a SS Gaussian, and continuous SS models, which for each $t \in \mathbb{Z}$ model in discrete-time [36]. We focus on (possibly) non-linear,

A. State Space Model

III. SYSTEM MODEL AND PRELIMINARIES

$\{\cdot\}$, $\|\cdot\|$, and $\mathbb{E}[\cdot]$, respectively. The Gaussian distribution with mean u and covariance Σ is denoted by $N(u, \Sigma)$. Finally, \mathbb{R} and \mathbb{Z} are the sets of real and integer numbers, respectively.

concluding remarks and future work.

The rest of this paper is organized as follows: Section II reviews the SS model and its associated tasks, and discusses related works. Section III details the proposed KalmannNet, Section IV presents the numerical study. Section V provides assumptions, as we discuss in detail below.

The proposed KalmannNet leverages data and partial domain knowledge to learn the filtering operation, rather than using data to explicitly estimate the missing SS model parameters. Although there is a large body of work that combines SS models with DNNs, e.g., [29]–[35], these approaches are sometimes used for different SS related tasks (e.g., smoothing, impatation); with a different focus, e.g., incorporating high-dimensional visual observations to a KE; or under different

4) We evaluate KalmannNet in various SS models. The experimental scenarios include synthetic setups, tracking the chaotic Lorenz system, and localization using the Michiganan NCLT data set [28]. KalmannNet is shown to converge much faster compared with purely DD systems, while outperforming the MB EKF, UKF, and PF, when using short trajectories.

Our main contributions are summarized as follows:

- 1) We design Kalmannet, which is an interpretable, low complexity, and data-efficient DNN-based real-time state estimator. Kalmannet builds upon the flow and theoretical principles of the KF, incorporating partial domain knowledge of the underyling SS model in its operation.
- 2) By learning the KG, Kalmannet circumvents the dependency of the KF on knowledge of the underlying noise statistics, thus bypassing numerically problematic matrix inversions involved in the KF equations and overcoming the need for tailored solutions for non-linear systems; e.g., filtering from data in a manner that is invariant to the sequence length. Specifically, we present an efficient supervised training scheme that enables Kalmannet to operate with arbitrary long trajectories while only training 3) We show that Kalmannet learns to carry out Kalmann approximation to handle non-linearities as in the EKF.

classic KF, and the model-agnostic nature of DNNs. Therefore, we build upon the success of our previous work in MB deep learning for signal processing and digital communication applications [24]–[27] to propose a hybrid MB/DD online recursive filter, coined KalmanNet. In particular, we focus on real-time state estimation for continuous-value SS models for which the KF and its variants are designed. We assume that the noise statistics are unknown and the underlying SS model is partially known or approximated from a physical model of the system dynamics. To design KalmanNet, we identify the KF as a critical component encapsulating the dependency on noise statistics and domain knowledge, and replace it with a compact RNN of limited complexity that is integrated into the KF flow. The resulting system uses labeled data to learn to carry out Kalman filtering in a supervised manner.

A common strategy when using DNNs is to encode the observations into some latent space that is assumed to obey a simple SS model, typically a linear Gaussian one, and track the state in the latent domain as in [56], [60], [61], or to use DNNs to estimate the parameters of the SS model as in [62], [63]. Tracking in the latent space can also be extended by applying a DNN decoder to the estimated state to return to the observations domain, while training the overall system end-to-end [31], [64]. The latter allows to design trainable systems for recovering missing observations and predicting future ones by assuming that the temporal relationship can be captured as an SS model in the latent space. This form of DNN-aided systems is typically designed for unknown or highly complex SS models, while we focus in this work on setups with partial domain knowledge, as detailed in Subsection II-B.

When the underlying dynamics of a system's dynamics are complex and only partially known or the emission model is intricate and cannot be captured in a closed form—e.g., visual observations as in a computer vision task [56]—one can resort to approximations and to the use of DNNs. Variational inference [57]—[59] is commonly used in connection with SS models, as in [29]–[31], [33], [34], by casting the Bayesian inference task to an optimization of a parameterized posterior and maximizing it directly to state recovery in real-time, as we consider here, or by casting the task to state recovery in real-time, as we consider here, as an objective. Such approaches cannot typically be applied to approximate procedures that tends to be complex and prone to errors.

When one can bound the uncertainty in the SS model in advance, an alternative approach to learning is to minimize the worst-case estimation error among all expected SS models. Such robust variations were proposed for various state estimation algorithms, including Kalman variatants [15]–[17], [53] and particle filters [54], [55]. The fact that these approaches aim to design the filter to be suitable for multiple different SS models typically results in degraded performance compared to operating with known dynamics.

approach to deal with partially known SS models is to impose a parametric model and then estimate its parameters. This can be achieved by jointly learning the parameters and state sequence using expectation maximization [46]–[48] and Bayesian probabilistic algorithms [38], or by selecting from a set of *a priori* known models [49]. When training data is available, it is commonly used to tune the missing parameters in advance, in a supervised or an unsupervised manner, as done in [50]–[52]. The main drawback of these strategies is that they are restricted to an imposed parametric model on the underlying dynamics (e.g., Gaussian noises).

A key ingredient in recursive Bayesian filtering is the update operation; namely, the need to update the prior estimate using new observed information. For linear Gaussian SS using the KF, this boils down to computing the KG. While the KF assumes linear SS models, many problems encountered in practice are governed by non-linear dynamics, for which one should resort to approximations. Several extensions of the EKF were proposed to deal with non-linearities. The EKF [7], [8] is a quasi-linear algorithm based on an analytical linearization of the SS model. More recent non-linear variations are based on numerical integration: UKF [10], the Gauss-Hermite Quadrature [39], and the Cubature KF [40]. For more complex SS models, and when the noise cannot be modeled as Gaussian, multiple variants of the PF were proposed that are based on sequential MC [11]–[13], [41]–[45]. These MC algorithms are considered to be asymptotically exact but relatively computationally heavy when compared to Kalman-based algorithms. These MB algorithms require accurate knowledge of the SS model, and their performance degrades in the presence of model mismatch.

C. Related Work

- Knowledge of the noise distribution of the noise signals w_t , and v_t , is not available.
 - The function $f(\cdot)$ and $h(\cdot)$ may constitute an approximation of the true underlying dynamics. Such approximations can correspond, for instance, to the representation of continuous time dynamics in discrete time, acquisition of continuous time dynamics in discrete time, and other forms of mismatches. While we focus on filtering in partially known SS models, we assume that we have access to a labeled data set containing a sequence of observations and their corresponding ground truth states. In various scenarios of interest, one can assume sensors both internally or externally to collect the ground truth truth data using offline and more computationally intensive algorithms. Finally, the inference complexity of the learned filter should be of the same order (and preferably smaller) as that of MB filters, such as the EKF.

x_t , based on each incoming observation y_t , in an online manner. Our main focus is on scenarios where one has *partial knowledge* of the SS model that describes the underlying dynamics. Namely, we know (or have an approximation of) the state-evolution function (*transition*) function $f(\cdot)$ and the state-observation emission function $h(\cdot)$. For real world applications, this knowledge is derived from our understanding of the system dynamics, its physical design, and the model of the sensors. As opposed to the classical assumptions in KF, the noise statistics of R and K are not known. More specifically, we assume:

We formulate KalmannNet by identifying the specific components of the EKF that are based on unavailable knowledge. As detailed in Subsection II-B, the functions $f(\cdot)$ and $h(\cdot)$ are known (though perhaps inaccurately); yet the covariance matrices Q and R are unavailable. These missing statistical moments are used in MB Kalmann filter only for computing the KG (see Fig. I). Thus, we design KalmannNet to learn the KG from data, and combine the learned KG in the overall EKF. In each time instance $t \in \mathcal{T}$, similarly to the EKF, KalmanNet estimates \hat{x}_t in two steps; prediction and update. In particular, a prior estimate for the current state \hat{x}_{t-1} is computed from \hat{x}_{t-1} via (8a). As opposed to its MB counterpart, KalmannNet does not rely on the knowledge of noise distribution and does not maintain an explicit estimate of the second-order statistical moments.

- 1) The prediction step is the same as in the MB EKF, except that only the first-order statistical moments are predicted.
- 2) The prediction step is the same as in the EKF, except that only the first-order statistical moments are predicted.

A. High Level Architecture

Here, we present KalmannNet; a hybrid, interpretable, data-efficient architecture for real-time state estimation in non-linear dynamical systems with partial domain knowledge. KalmannNet combines MB Kalmann filtering with an RNN to cope with model mismatch and non-linearity. To introduce KalmannNet, we begin by explaining its high-level operation in Subsection III-A. Then we present the features processed by its internal RNN and the specific architectures considered for implementation and training in Subsections III-B-III-D. Finally, we provide a discussion in Subsection III-E.

III. KALMANN ET

Fig. 1: EKF block diagram. Here, Z_{-1} is the unit delay. An illustration of the EKF is depicted in Fig. 1. The resulting filter admits an efficient linear recursive structure. However, it requires full knowledge of the underlying model and notably degrades in the presence of model mismatch. When the model is highly non-linear, the local linearity approach may not hold, and the EKF can result in degraded performance. This motivates the augmentation of the EKF into the deep learning-aided KalmanNet, detailed next.

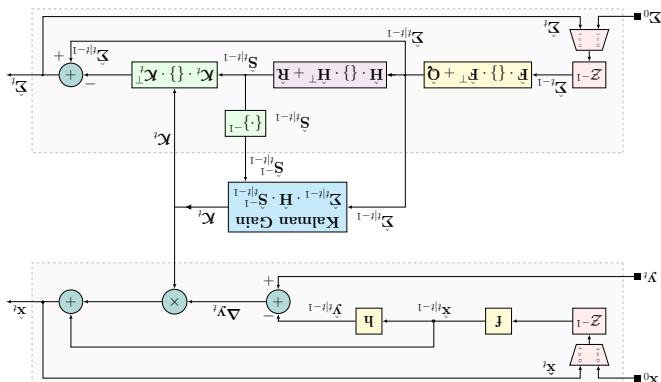


Fig. 1: EKF block diagram. Here, Z_{-1} is the unit delay.

The EKF extends the KF for non-linear $f(\cdot)$ and/or $h(\cdot)$, as in (1). Here, the first-order statistical moments (3a) and (4a) are replaced with $\hat{x}_{t|t-1}$. The second-order moments, though, cannot be propagated through the non-linearity, and must thus be approximated. The EKF linearizes the differenceable $f(\cdot)$ and $h(\cdot)$ in a time-dependent manner using their partial derivative matrices, also known as Jacobians, evaluated at $\hat{x}_{t|t-1}$ and where F_t is plugged into (3b) and H_t is used in (4b) and (6). When the SS model is linear, the EKF coincides with the KF, which achieves the MSE for linear Gaussian SS models.

(9a)
$$\hat{F}_t = \mathcal{L}^f(\hat{x}_{t-1|t-1})$$

(9b)
$$\hat{H}_t = \mathcal{L}^h(\hat{x}_{t-1|t-1})$$

(11)
$$\Delta y_t = y_t - \hat{y}_{t|t-1}.$$

$$\cdot \left(1 - t | \tau_{\mathbf{X}} \right) \mathbf{u}_{\mathcal{L}} = t \mathbf{H}$$

$$(9a) \quad \mathbf{H}^t \equiv \int_{t-1}^t \mathbf{x}^{t-t}$$

\hat{x}_{t-1} . Namely,

respectively. The second-order moments, though, cannot be propagated through the non-linearity, and must thus be approximated. The EKF linearizes the differentiable $f(\cdot)$ and proximates, also known as Jacobians, evaluated at x_{t-1}^{t-1} and $h(\cdot)$ in a time-dependent manner using their partial derivative matrices, also known as Jacobians, evaluated at x_{t-1}^{t-1} and

$$\cdot \left(\begin{smallmatrix} 1-t & t \\ & x \end{smallmatrix} \right) h = \begin{smallmatrix} 1-t & t \\ & x \end{smallmatrix} Y$$

$$(8a) \quad \mathbf{x}^{t-1} = \mathbf{f}(\mathbf{x}^{t-1}),$$

are replaced with

The EKF extends the KF for non-linear $f(\cdot)$ and/or $h(\cdot)$, as in [1]. Here, the first-order statistical moments ($\hat{3}a$) and ($\hat{4}a$)

$$\cdot \tau^{-\tau} | \tau \mathcal{A} - \tau \mathcal{A} = \tau \mathcal{A} \nabla$$

The term ΔY_t is the innovation; i.e.,

$$\Delta Y_t = \sum_{j=1}^{t-1} S_j \cdot H^{-1} \cdot e_j$$

 between the predicted observation and the
 and it is the only term that depends on the

$$\cdot \tau^{-\ell} |_{\mathcal{V}} \mathbf{s} : H \cdot \tau^{-\ell} |_{\mathcal{V}} \gamma = \tau \gamma$$

$$\cdot {}^t\mathcal{K} \cdot {}^{1-t}\mathcal{S} \cdot {}^t\mathcal{K} - {}^{1-t}\mathcal{S} = {}^t\mathcal{K}$$

$$(\mathfrak{g}_C) \quad \quad \quad {}^7\mathcal{F}\nabla : {}^7\mathbf{v} + {}^{1-t}{}^7t\mathbf{x} \equiv {}^7{}^7\mathbf{x}$$

2) In the update step, the *a posteriori* state moments are computed based on the *priori* moments as

$$(4b) \quad H \cdot \sum_{t=1}^{|t|-1} H + R.$$

$$(a) \quad \mathbf{I} - t^2 \mathbf{x} \cdot \mathbf{H} = \mathbf{I} - t^2 \mathbf{y}$$

1) The first step *practices* the current *a priori* statistical moments based on the previous *posterior* estimates. Specifically, the moments of \mathbf{x} are computed using the knowledge of the evolution matrix \mathbf{F} as

$$\mathbf{x}_{t|t-1} = \mathbf{F} \cdot \mathbf{x}_{t-1|t-1}, \quad (3a)$$

$$\mathbf{z}_{t|t-1} = \mathbf{F} \cdot \mathbf{z}_{t-1|t-1} + \mathbf{Q} \quad (3b)$$

and the moments of the observations \mathbf{y} are computed based on the knowledge of the observation matrix \mathbf{H} as

It is extended into the EKF for non-linear SS models.

Our proposed KalmannNet, detailed in the following section, is based on the MB-KF, which is a linear recursive estimator. In every time step t , the KF produces a new estimate \hat{x}_t using only the previous estimate \hat{x}_{t-1} as a sufficient statistic and the new observation y_t . As a result, the computational complexity of the KF does not grow in time. We first describe the original algorithm for linear SS models, as in (2), and then discuss how

D. Model-Based Kalman Filtering

V1 Direct application of BPTT, where for each training iteration, gradients are computed over the entire trajectory.

V2 An application of the truncated BPTT algorithm [70]. Here, given a data set of long trajectories (e.g., $T = 3000$ time steps), each long trajectory is divided into multiple short trajectories (e.g., $T = 100$ time steps), which are shuffled and used during training.

V3 An alternative application of truncated BPTT, where we

Since KalmannNet is a recursive architecture with both extreme recurrence and an internal RNN, we use the backpropagation through time (BPTT) algorithm [69] to train KalmannNet across time. Specifically, we unfold KalmannNet across time with shared network parameters, and then compute a forward and backward gradient estimation pass through the network. We consider three different variations of applying the BPTT algorithm:

$$\mathcal{L}_k(\Theta) = \frac{1}{M} \sum_{j=1}^M \ell^{i_k}_{y_k}(\Theta). \quad (14)$$

To optimize Θ , we use a variant of mini-batch stochastic gradient descent in which for every batch indexed by k , we choose $M < N$ trajectories indexed by i_1^k, \dots, i_M^k , computing the mini-batch loss as

$$(13) \quad \|\Theta\|_2 + \frac{1}{L} \left\| \left(\Theta - (\Theta - \frac{1}{L} X)^T \right)^T X \right\|_2 = (\Theta - \frac{1}{L} X)^T X$$

By letting Θ denote the trainable parameters of the RNN, we then construct an ℓ_2 -regularized mean-squared error (MSE) loss measure and γ be a regularization coefficient, we have

$$\mathbf{A}_i = \left[\mathbf{y}_{(i)}^1, \dots, \mathbf{y}_{(i)}^{T_i} \right], \quad \mathbf{X}_i = \left[\mathbf{x}_{(i)}^0, \dots, \mathbf{x}_{(i)}^{T_i} \right]. \quad (12)$$

– *if we are missing necessary information, we can use a guess*

The data set used for training comprises N trajectories that can be of varying lengths. Namely, by letting T_i , be the length of the i -th training trajectory. The data set is given by $D = \{x_i\}_{i=1}^N$

the KG for training purposes.

where $\Delta x_t \triangleq x_t - \hat{x}_{t-t-1}$. The gradient computation in (11) indicates that one can learn the computation of the KG by training KalmannNet end-to-end using the squared-error loss.

$$\frac{\partial \mathcal{A}^t}{\partial \mathbf{x}^t} = \frac{\partial \mathcal{A}}{\partial \mathbf{x}^t} - \frac{\partial \mathcal{A}^t}{\partial \mathbf{x}^t} \nabla_{\mathbf{x}^t} \mathcal{A}^t$$

which is also used to evaluate the MB KF. By doing so, we build upon the ability to backpropagate the loss to the computation of the KG. One can obtain the loss gradient with respect to the KG from the output of KalmanNet since

$$\| \mathbf{\hat{x}} - \mathbf{x} \|_2 = \mathcal{J} \quad (10)$$

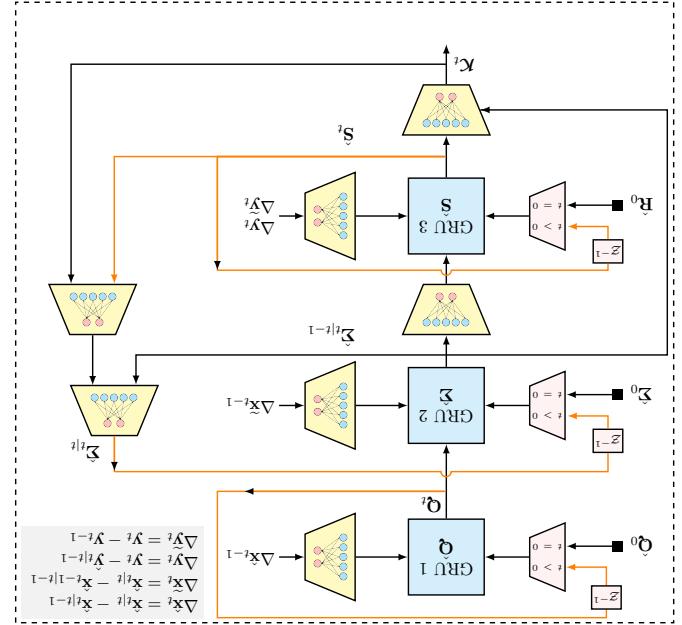
estimates $x_t^{[t]}$, we train KalmanNet end-to-end. Namely, we compute the loss function \mathcal{L} based on the state estimate \hat{x}_t , which is not the output of the internal RNN. Since this vector takes values in a continuous set \mathbb{R}^m , we use the squared-error loss

KlamNet is trained using the available labeled data set in a supervised manner. While we use a neural network for computing the KG rather than for directly producing the

D. Training Algorithm

The second architecture uses separate GRU cells for each of the tracked second-order statistical moments. The division of the architecture into separate GRU cells and FC layers and their interconnection is illustrated in Fig. 4. As shown in the figure, the network comprises three GRU layers, connected in a cascade with dedicated input and output FC layers. The first GRU layer tracks the unknown state noise covariance Q , thus tracking m_2 variables. Similarly, the second and third GRUs track the predicted moments Z^{t+1} (3b) and S^t , thus having m_2 and n_2 hidden state variables, respectively. The GRUs are interconnected such that the learned Q is used to compute Z^{t+1} , which in turn is used to obtain S^t , while both Z^{t+1} and S^t are involved in producing K^t . This architecture is composed of a non-standard interconnection between GRUs and FC layers, is more directly tailored towards the formulation of the SS model and the operation of the MBKF compared with the simpler first architecture. As such, it provides lesser abstraction; i.e., it is expected to be more less reliable parameters. For instance, in the numerical study reported in Subsection IV-D, utilizing the first architecture requires the order of $5 \cdot 10^5$ trainable parameters, while the second architecture utilizes merely $2 \cdot 5 \cdot 10^4$ parameters.

Fig. 4: Kalmannet RNN block diagram (architecture #2). The input features are used to update three GRUs with dedicated FC layers, and the overall interconnection between the blocks is based on the flow of the KG computation in the MB KE.



Finally, we note that while we focus here on filtering with finite-variance communication channels, SS models are used to represent additional related tasks, SS problems such as smoothing and prediction, as discussed in Subsection II-A. The fact that KalmanNet does not explicitly estimate the SS model implies that it cannot simply substitute these parameters into an alternative algorithm capable of carrying out tasks other than filtering. Nonetheless, one can still design DNN-aided algorithms for these tasks operating with partially known SS models as extensions of KalmanNet.

Our design of Kalmannet gives rise to many interesting future extensions. Since we focus here on SS models where the mappings $f(\cdot)$ and $h(\cdot)$ are known up to some approximation errors, a natural extension of Kalmannet is to use the data to pre-estimate them, as demonstrated briefly in the numerical study. Another alternative to cope with these approximation errors is to utilize dedicated neural networks to learn these mappings while training the entire model in an end-to-end fashion. Doing so is expected to allow Kalmannet to be utilized in scenarios with analytically intractable SS models, as often arises when tracking based on unsstructured observations.

RNNs for end-to-end state estimation, and also approaches the MSE performance achieved by the MB KF in linear Gaussian SS models. Furthermore, the fact that KalmanNet preserves the flow of the EKF implies that the intermediate features exchanged between its modules have a specific operation meaning, providing interpretability that is often scarce in end-to-end, deep learning systems. Finally, the fact that KalmanNet learns to compute the KG indicates the possibility of providing not only estimates of the state x^t , but also a measure of confidence in this estimate, as the KG can be related to the covariance of the estimate, as initially explored in [7].

Furthermore, KalmanNet is derived for SS models when noise statistics are not specified explicitly. A MB approach to tackle this without relying on data employs the robust Kalman filter [15]–[17], which designs the filter to minimize the maximal MSE within some range of assumed SS models, at the cost of performance loss, compared to the true model. When one has access to data, the direct strategy to implement the EKF in such setups is to use the data to estimate Q and R , either directly from the data or by backpropagating through the operation of the EKF as in [51], and utilize these estimates to compute the KG. As covariance estimation can be a challenging task when dealing with high-dimensional signals, KalmanNet bypasses this need by directly learning the KG, and by doing so approaches the MSE of MB Kalman filtering with full knowledge of the SS model, as demonstrated in Section VI. Finally, the computation complexity for each time step $t \in T$ is also linear in the RNN dimensions and does not involve matrix inversion. This implies that KalmanNet is a good candidate to apply for high dimensional SS models and follows the flow of MB Kalman filtering rather than being effs from its model awareness and the fact that its operation is limited to purely DD state estimation, KalmanNet benefits from its from its model awareness and the fact that its operation is limited to purely DD state estimation, KalmanNet is able to utilize a black box. As numerically observed in Section IV, KalmanNet achieves improved MSE compared to utilizing a black box.

E. Discussion

Overall, directly applying BPTT via V_1 may be computationally expensive and unstable. Therefore, a favored approach is to first use the truncated BPTT as in V_2 as a warm-up phase (train first on short trajectories) in order to stabilize its learning process, after which Kalmannet is tuned using V_3 . The procedure in V_3 is most suitable for systems that are known to be likely to quickly converge to a steady state (e.g., linear SS models). In our numerical study, reported in Section IV, we utilize all three approaches.

2) *Neural Model Selection*: Next, we evaluate and confirm our design and architectural choices by considering a 2 × 2

1) *Full Information*: We start by comparing KalmanNet of setting CI to the MB KF for the case of full information, where the latter is known to minimize the MSE. Here, we set H to take the inverse camouflaged form, and $v = 0$ [dB]. To demonstrate the applicability of KalmanNet to various linear systems, we experiment with systems of different dimensions; namely, $m \times n \in \{2 \times 2, 5 \times 5, 10 \times 1\}$, and with transfer functions; namely, $T \in \{50, 100, 150, 200\}$. In Fig. 3a we can clearly observe that KalmanNet achieves the MSE of the MB KF. Moreover, to further evaluate the gains of the hybrid architecture of KalmanNet, we check that its learning is transferable. Namely, in some of the experiments, we test KalmanNet on longer trajectory initial conditions. The fact that it was trained on, and with different initial conditions, shows that it was able to learn the system dynamics and adapt to the new conditions.

Our first experimental study compares KalmanNet to the MBKF for different forms of system dynamics. Unless stated otherwise, here F takes the form of a linear state-space model. KalmanNet generates a linear state-space model of the system dynamics, while MBKF generates a non-linear state-space model. The two models are compared by their ability to predict the system's future state given its current state and past inputs.

B. Linear State Space Model

2) **Model-Based Filters:** In the following experiments we compare KalmanNet with several MB filters. For study we used the software package [75], while the PF is implemented based on [76] using 100 particles and without parallelization. During our numerical study, when model uncertainty was introduced, we optimized the performance of the MB algorithms by carefully tuning the covariance matrices, usually via a grid search. For long trajectories (e.g., $T > 1500$) it was sometimes necessary to tune these matrices, even in the case of full information, to compensate for inaccurate uncertainty propagation due to non-linear approximations and to avoid divergence.

- CI KalmannNet architecture #1 with input features {F₂, F₄} and with training algorithm V₃.
- CII KalmannNet architecture #1 with input features {F₂, F₄} and with training algorithm V₁.
- CIII KalmannNet architecture #1 with input features {F₂, F₄} and with training algorithm V₂.
- CIV KalmannNet architecture #1 with input features {F₁, F₃, F₄} and with training algorithm V₂.
- CV KalmannNet architecture #2 with all input features and with training algorithm V₁.
- CVI KalmannNet architecture #2 with all input features and with training algorithm V₂.
- CVII KalmannNet architecture #2 with all input features and with training algorithm V₁.
- CVIII KalmannNet was trained using the Adam optimizer [4].

- 1) **KalmannNet Settings:** In Section III we present several architectures and training mechanisms that can be used when implementing KalmannNet. In our experimental study we consider three different configurations of KalmannNet:

In some of our experiments, we evaluate both the MSE and its standard deviation, where we denote these measures by f and \hat{f} , respectively.

The source code is released in the [SNP](https://github.com/KalamaniT/SNP) repository along with the complete set of preparation scripts used in the SNP metric evaluation can be found online at <https://github.com/KalamaniT/SNP>.

By (15), setting ν to be 0 DB implies that both the state noise and the observation noise have the same variance. For consistency, we use the term *full information* for cases where the SS model available to KalmanNet and its MB counterparts accurately represents the underlying dynamics. More specifically, KalmanNet operates with full knowledge of $f(\cdot)$ and $h(\cdot)$, and without access to the noise covariance of $\mathbf{f}(\cdot)$. More specifically, KalmanNet and its MB counterparts operate with full knowledge of \mathbf{Q} and \mathbf{R} . The term *partial information* refers to the case where KalmanNet and its MB counterparts operate with some level of model mismatch, where the SS model design parameters do not represent the SS model metric used to evaluate the performance is the MSE on a [dB] scale. In the figures we depict the MSE in [dB] versus the scale. In the figures we depict the MSE in [dB] versus the scale. In the figures we depict the MSE in [dB] versus the scale. In the figures we depict the MSE in [dB] versus the scale.

$$\mathbf{Q} = \mathbf{b}_2 \cdot \mathbf{I}, \quad \mathbf{R} = \mathbf{i}_2 \cdot \mathbf{I}, \quad \text{and} \quad (15)$$

Throughout the numerical study and unless stated otherwise,
in the experiments involving synthetic data, the SS model is
generated using diagonal noise covariance matrices; i.e.,

A. Experimental Setting

(c) In our last study we consider a *localization* use case based on the Michigean NCLT data set [28]. Here, we compare KalmannNet with MBKF that assumes a linear Wiener kinematic model [36] and with a *vanilla* RNN based end-to-end state estimator, and demonstrate the ability of KalmannNet to track real world dynamics that was not synthetically generated from an underlying SS model.

(a) In our first experimental study, we consider multiple linear SS models, and compare Kalmannet to the MBKF which is known to minimize the MSE in such a setup. We also confirm our design and architectural choices by comparing it to various benchmark algorithms:

In this section we present an extensive numerical study of Kalmannet, evaluating its performance in multiple setups and comparing it to various benchmark algorithms:

(b) In our first experimental study, we consider multiple linear SS models, and compare Kalmannet to the MBKF which is known to minimize the MSE in such a setup. We also confirm our design and architectural choices by comparing it to various benchmark algorithms:

IV. EXPERIMENTS AND RESULTS

KF. For instance, as the MB KF constitutes the first part of the Rauch-Tung-Streibel smoother [73], one can extend Kalman-Net to implement high-performance smoothing in partially known SS models, as we have recently begun investigating in [67]. Nonetheless, we leave the exploration of extensions of KalmanNet to alternative tasks associated with SS models for future work.

Such scenarios represent a setup in which the analytical approachimation of the SS model differs from the true generative model. The resulting MSE curves depicted in Fig. 6a demonstrate that KalmannNet (with setting C) achieves a gain over the MBKF. In particular, despite the fact that KalmannNet overcomes a minor gap from the true generative model, it learns to apply an alternative KG, resulting in MSE within a minor gap from the MSE; i.e., from the KF with the true state-observation mismatch while setting $T = 100$ and $\nu = -20$ [dB]. The model mismatch is achieved by using a rotated observation matrix $H^{x=10}$ for data generation, while sensors exists. The resulting achieved MSE depicted in Fig. 6b represents a setup in which a slight misalignment ($\approx 5\%$) of the sensors' locations with respect to the observation matrix $H = I$ as the observation design. Such scenarios demonstrate that KalmannNet (with setting C) converges to within a minor gap from the MSE. Here, we performed an additional experiment, first estimating the observation matrix from data, and then KalmannNet used the estimate matrix denoted H_a . In this case it is observed in Fig. 6b that KalmannNet achieves the MSE lower bound. These results imply that KalmannNet converges also in distribution to the true generative model.

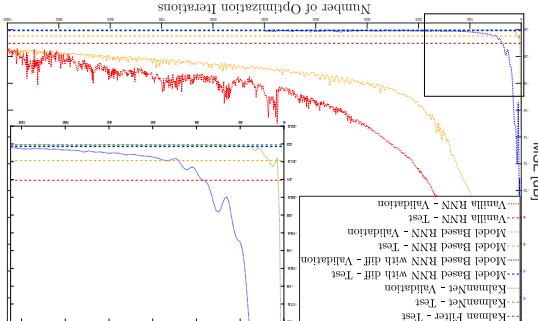
$$\mathbf{F}^{\alpha_0} = \mathbf{R}_{xy}^{\alpha_0} \cdot \mathbf{F}_0, \quad \mathbf{R}_{xy}^{\alpha_0} = \begin{pmatrix} \cos \alpha & -\sin \alpha \\ \sin \alpha & \cos \alpha \end{pmatrix}. \quad (16)$$

Model mismatch, we have compared the results of numerical experiments with the state-space model (H) or in the state-observation model (F). We simulate a 2×2 SS model with mismatches in either the state-evolution model mismatch as a result of partial model information. We model mismatch as a result of partial model information. We simulate a 2×2 SS model with mismatches in either the state-evolution model (F) or in the state-observation model (H).

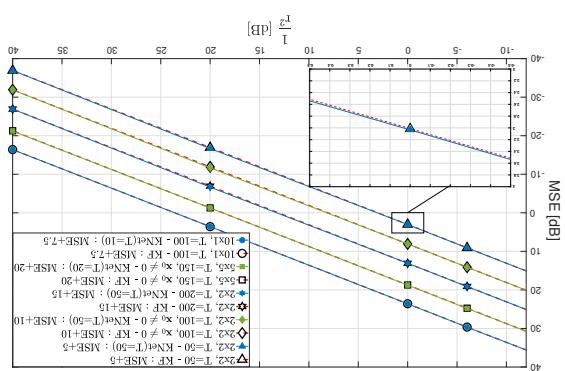
TABLE I: Test MSE in [dB] when trained using $L = 20$.

Test T	Vanilla RNN	MBRNN	MBRNN	KalmanNet	KF
20	-20.98	-21.53	-21.92	-21.92	-21.97
200	58.14	36.8	-21.88	-21.90	-21.91

(b) Learning curves for DD state estimation.



1.5-2.1



- To further evaluate the gains of KalmannNet over end-to-end RNNs, we compare the pre-trained models using trajectories with different initial conditions and a longer time horizon ($T = 200$) than the one on which they were trained ($T = 20$). The results, summarized in Table I, show that KalmannNet maintains achieving the MSE, as already observed in Fig. 5a. The MB RNN and vanilla RNN are more than 50 [dB] from the MSE, implying that their learning is not transferable and that they do not learn to implement Kalmann filtering. However, a minor gap of that achieved by the MB RNN. The results of this study validate the considerations used in designing KalmannNet for the DD filtering problem discussed in Subsection II-B.

- *Vanilla RNN* directly maps the observed y_t to an estimate of the state x_t .
 - *MB RNN* imitates the Kalmann filtering operation by first recovering x_{t-1} using domain knowledge, i.e., via (3a), and then uses the RNN to estimate an increment Δx_t from the prior to posterior.
 - All RNNs utilize the same architecture as in KalmannNet with a single GRU layer and the same learning hyperparameters. In this experiment we test the trained models on trajectories with the same length as they were trained on, namely $T = 20$. We can clearly observe how each of the key design considerations of KalmannNet affect the learning curves depicted in Fig. 5b:
 - The incorporation of the known SS model allows the MB RNN to outperform the vanilla RNN, although both converge slowly and fail to achieve the MSE.
 - Using the sequences of differences as input notably improves the convergence rate of the MB RNN, indicating the benefits of using the differences as features, as discussed in Subsection III-B.
 - Learning is further improved by using the RNN for recovering the state x_t rather than for directly estimating x_t .

via (20) with $j = 5$.

MB filter when using the state-evolution matrix \mathbf{F} computed 1) *Full Information*: We first compare KalmannNet to the

studies, as discussed in the sequel. The discrete-time state-evolution model uses for $\mathbf{F}(\cdot)$ and $\mathbf{h}(\cdot)$ varying between the different tently invariant of the distribution of the noise signals, with the general. In the following experiments, KalmannNet is considered with $\Delta\tau = 0.02$ sampling interval. In the Taylor order and $\Delta\tau = 5$ Taylor series is used for generating the simulated Lorenz attractor data. Unless stated otherwise the data was process noise, is used in (21), with additional discrete-time state-evolution model in (21), with additional

$$\mathbf{x}_{t+1} = \mathbf{f}(\mathbf{x}_t) = \mathbf{F}(\mathbf{x}_t) \cdot \mathbf{x}_t. \quad (21)$$

The resulting discrete-time evolution process is given by

$$\mathbf{F}(\mathbf{x}_t) \triangleq \exp(\mathbf{A}(\mathbf{x}_t) \cdot \Delta\tau) \approx \mathbf{I} + \sum_{j=1}^5 \frac{(\mathbf{A}(\mathbf{x}_t) \cdot \Delta\tau)^j}{j!}. \quad (20)$$

Finally, we take the Taylor series expansion of (19) and finite series approximation (with j coefficients), which results in

$$\mathbf{x}_{t+\Delta\tau} = \exp(\mathbf{A}(\mathbf{x}_t) \cdot \Delta\tau) \cdot \mathbf{x}_t. \quad (19)$$

interval $\Delta\tau$, is constant in a small neighborhood of \mathbf{x}_t ; i.e.,

$$\mathbf{A}(\mathbf{x}_t) \approx \mathbf{A}(\mathbf{x}_{t+\Delta\tau}).$$

To get a discrete-time, state-evolution model, we repeat the sampling interval $\Delta\tau$ and assume that $\mathbf{A}(\mathbf{x}_t)$ can be kept steps used in [35]. First, we sample the noiseless process with

$$\frac{\partial}{\partial t} \mathbf{x}_t = \mathbf{A}(\mathbf{x}_t) \cdot \mathbf{x}_t, \quad \mathbf{A}(\mathbf{x}_t) = \begin{pmatrix} 0 & x_{1,t} & -\frac{8}{3} \\ -10 & 10 & 0 \end{pmatrix}. \quad (18)$$

time process \mathbf{x}_t with $t \in \mathbb{R}_+$ is given by

In particular, the noiseless state-evolution of the continuous-

	$1/\epsilon^2$ [dB]	-12.04	-6.02	0	20	40
KalmannNet	$\frac{\partial}{\partial t}$	-7.25	-13.19	-19.22	-39.13	-59.10
PF	$\frac{\partial}{\partial t}$	-6.59	-13.33	-18.78	-26.70	-30.98
UKF	$\frac{\partial}{\partial t}$	-0.91	-1.34	-1.18	-2.24	-3.72
EKF	$\frac{\partial}{\partial t}$	-2.99	-0.63	-0.89	-0.45	-0.3
		-1.23	-0.57	-0.50	-0.55	-0.31
		-1.34	-0.47	-0.43	-0.43	-0.44
		-1.84	-0.49	-0.50	-0.50	-0.49
		-1.82	-0.53	-0.53	-0.53	-0.53
		-1.81	-0.50	-0.50	-0.50	-0.50
		-1.80	-0.47	-0.47	-0.47	-0.47
		-1.79	-0.43	-0.43	-0.43	-0.43
		-1.78	-0.40	-0.40	-0.40	-0.40
		-1.77	-0.37	-0.37	-0.37	-0.37
		-1.76	-0.34	-0.34	-0.34	-0.34
		-1.75	-0.31	-0.31	-0.31	-0.31
		-1.74	-0.28	-0.28	-0.28	-0.28
		-1.73	-0.25	-0.25	-0.25	-0.25
		-1.72	-0.22	-0.22	-0.22	-0.22
		-1.71	-0.19	-0.19	-0.19	-0.19
		-1.70	-0.16	-0.16	-0.16	-0.16
		-1.69	-0.13	-0.13	-0.13	-0.13
		-1.68	-0.10	-0.10	-0.10	-0.10
		-1.67	-0.07	-0.07	-0.07	-0.07
		-1.66	-0.04	-0.04	-0.04	-0.04
		-1.65	-0.01	-0.01	-0.01	-0.01
		-1.64	0.02	0.02	0.02	0.02

TABLE IV: MSE [dB] – Synthetic non-linear SS model; partial

	$1/\epsilon^2$ [dB]	-12.04	-6.02	0	20	40
KalmannNet	$\frac{\partial}{\partial t}$	-7.25	-13.19	-19.22	-39.13	-59.10
PF	$\frac{\partial}{\partial t}$	-6.59	-13.33	-18.78	-26.70	-30.98
UKF	$\frac{\partial}{\partial t}$	-0.91	-1.34	-1.18	-2.24	-3.72
EKF	$\frac{\partial}{\partial t}$	-2.99	-0.63	-0.89	-0.45	-0.3
		-1.23	-0.57	-0.50	-0.55	-0.31
		-1.34	-0.47	-0.43	-0.43	-0.44
		-1.84	-0.49	-0.50	-0.50	-0.49
		-1.82	-0.40	-0.40	-0.40	-0.40
		-1.81	-0.37	-0.37	-0.37	-0.37
		-1.80	-0.34	-0.34	-0.34	-0.34
		-1.79	-0.31	-0.31	-0.31	-0.31
		-1.78	-0.28	-0.28	-0.28	-0.28
		-1.77	-0.25	-0.25	-0.25	-0.25
		-1.76	-0.22	-0.22	-0.22	-0.22
		-1.75	-0.19	-0.19	-0.19	-0.19
		-1.74	-0.16	-0.16	-0.16	-0.16
		-1.73	-0.13	-0.13	-0.13	-0.13
		-1.72	-0.10	-0.10	-0.10	-0.10
		-1.71	-0.07	-0.07	-0.07	-0.07
		-1.70	-0.04	-0.04	-0.04	-0.04
		-1.69	-0.01	-0.01	-0.01	-0.01
		-1.68	0.02	0.02	0.02	0.02

TABLE III: MSE [dB] – Synthetic non-linear SS model; full

mismatches due to sampling a continuous-time signal into trajectory and a real world practical challenge of handling continuous-time system of ordinary differential equations in continuous-time. This synthetic generated system solution to the Lorenz attractor is a three-dimensional chaotic

D. Lorenz Attractor

operation and low complexity of the KF. The MB variations of the KF fail, KalmannNet dynamics, where MB variation of the KF achieves better approximation of the underlying when using full information for such set-ups. We thus conclude that in the presence of harsh non-linearities as well as model uncertainty due to harsh noise levels; i.e., for $\frac{1}{\epsilon} = -12.04$ and its performance is within a small gap of that achieved in all experiments, KalmannNet overcomes such mismatches. In all experiments, KalmannNet achieves better performance of the MB filters due to the model mismatch. The true model, resulting in a notable degradation from evolution parameters used by the filters differs slightly from the presence of partial model information, the state-achieves superior MSE.

the KG from data, KalmannNet manages to overcome this and to a non-linear effect. Nonetheless, by learning to compute [dB], the MB EKF suffers from degraded performance due to noise levels; i.e., for $\frac{1}{\epsilon} = -12.04$ achieves the lowest MSE in such set-ups, and KalmannNet achieves similar MSE in the case of partial information and in the low noise regime, EKF achieves the lowest MSE values due to its ability to approach studies. For full information and in our main MB benchmark in our experimental fore serving as the lowest MSE values among the MB filters, there, achieves the lowest MSE values among the MB filters, the case of partial information. We first observe that the EKF in Table III for the case of full information, and in Table IV for the full evaluation with the MB EKF, UKF, and PF is given. The full evaluation with the MB EKF, UKF, and PF is given depicted in Fig. 7 for both full and partial model information. The MSE values for different levels of observation noise with setting C4.

In the following we generate trajectories of $T = 100$ time steps from the noisy SS model in (1), with $\nu = -20$ [dB], while using $\mathbf{f}(\cdot)$ and $\mathbf{h}(\cdot)$ as in (17) computed in a component-wise manner, with parameters as in Table II. KalmannNet is used

TABLE IV: MSE [dB] – Synthetic non-linear SS model; partial

TABLE III: MSE [dB] – Synthetic non-linear SS model; full

TABLE II: Non-linear toy problem parameters.

Partial	1	1	0	0	1	1	0
Full	0.9	1.1	0.1 π	0.01	1	b	c

SS model is given by

observation model is a second order polynomial. The resulting

evolution model takes a sinusoidal form, while the state-

Next, we consider a non-linear SS model, where the state-

C. Synthetic Non-Linear Model

KE.

State-of-observation rotation mismatch: Here, the presence of mismatch in the observations model is simulated by using data generated by an identity matrix rotated by merely $\theta = 1^\circ$. This rotation is equivalent to sensor misalignment of $\approx 0.55\%$. The results depicted in Figure 9b and reported in Table VII clearly demonstrate that this seemingly minor rotation can

State-evolution mismatch: In this study, both KalmanNet and the MB algorithms operate with a crude approximation of the evolution dynamics obtained by compounding (20) with $J = 2$, while the data is generated with an order $J = 5$ Taylor series expansion. We again set \mathbf{h} to be the identity mapping (20) and the evolution mismatch \mathcal{C}_4 learns to partially overcome this model mismatch, and reported in Table VI, demonstrates that KalmanNet with $T = 2000$, and $\nu = -20$ [dB]. The results, depicted in Fig. 9a and reported in Table VII, demonstrate that KalmanNet outperforms its MB counterparts operating with the same level of partial information.

Since the EKF produced the best results in the full information case among all non-linear MB filtering algorithms, we use it as a baseline for the MSE lower bound.

- State-evolution mismatch due to use of a Taylor series approximation of insufficient order.
 - State-observation mismatch as a result of misalignment due to rotation.
 - State-observation mismatch as a result of sampling from continuous-time to discrete-time.

2) *Partial Information:* We proceed to evaluate Kalmannet and compare it to its MB counterparts under partial model information. We consider three possible sources of model mismatch arising in the Lorenz attractor setup:

in such non-linear setups, the sub-optimal MB approaches operating with full information of the SS model are substantially outperformed by KalmanNet (with setting C4).

TABLE VI: MSE [dB] - Lorenz attractor with non-linear

Noisy non-linear observations: Next, we consider the case where the observations are given by a non-linear function of the current state, setting \mathbf{h} to take the form of a transformation from a cartesian coordinate system to spherical coordinates. We further set $T = 20$ and $\nu = 0$ [dB]. From Figure 8b and reported in Table VI we observe that depicted in Fig. 8b and reported in Table VI we observe that

Noisy state observations: Here, we set $h(\cdot)$ to be the identity transformation, such that the observations are noisy versions of the true state. Further, we set $\nu = -20 \text{ [dB]}$ and $T = 2000$. As observed in Fig. 8a, despite being trained on short trajectories $T = 100$, KalmannNet (with setting C3) achieves excellent MSE performance—namely, comparable to EKF—and outperforms the UKF and PF. The full details of the experiments are given in Table V. All the MB algorithms were optimized for performance; e.g., applying the EKF with full model information achieves an unstable state tracking performance, with MSE values surpassing 30 [dB]. To stabilize the EKF, we had to perform a grid search using the available data set to optimize the process noise Q used by the filter.

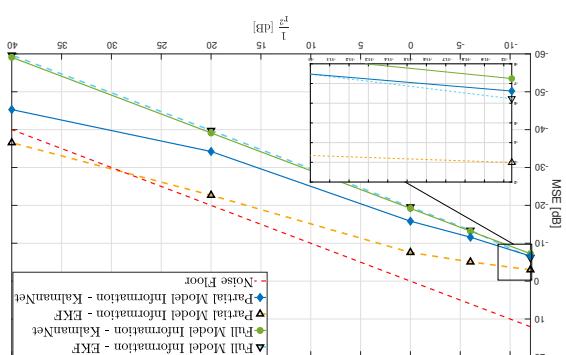
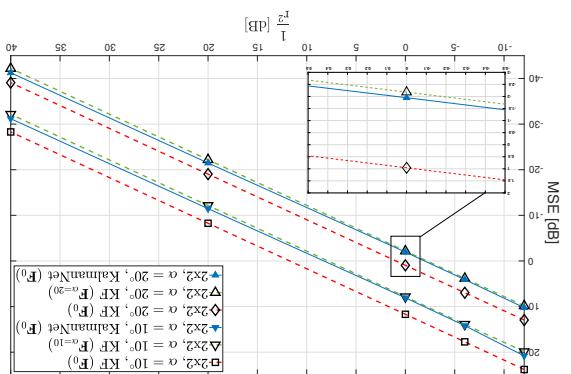
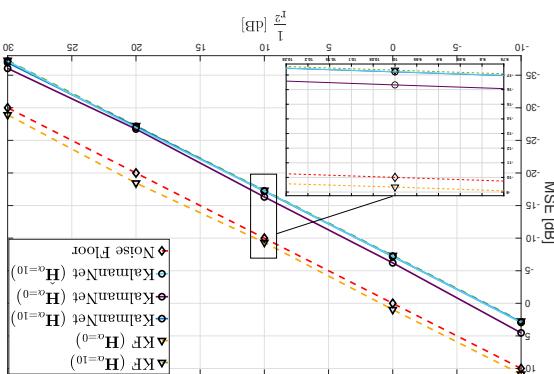


Fig. 6: Linear SS model, partial information.

(a) State-evolution mismatch.



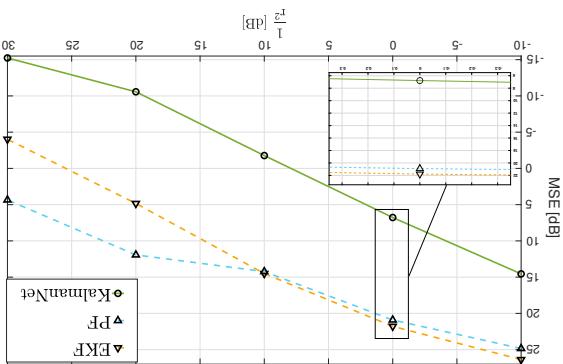
In our final experiment we evaluate KalmannNet on the Michiganan NCLT data set [28]. This data set comprises different labeled trajectories, with each one containing noisy sensor

E. Real World Dynamics: Michigan NCL Data Set

The resulting MSE values for $\frac{z}{\sigma} = 0$ [dB] of Kalmannet with configuration **C4** compared with the MB filters and with the end-to-end neural network trained MB-RNN (see Subsection IV-B) are reported in Table IX. The results demonstrate that Kalmannet overcomes the mismatch induced by representing a continuous-time SS model in discrete-time, achieving a substantial processing gain over the MB alternatives due to its learning capabilities. The results also demonstrate that Kalmannet significantly outperforms a straight-forward combination of domain knowledge; i.e. a state-transmission function $f(\cdot)$, with end-to-end RNNs. A fully model-agnostic RNN was shown to diverge when trained for this task. In Fig. 10 we visualize how this gain is translated into clearly improved tracking of a single trajectory. To show that these gains of Kalmannet do not come at the cost of computational slow-down the same platform – Google Colab with CPU: Intel(R) Xeon(R) CPU @ 2.20GHz, GPU: Tesla P100-Pcie-16GB. We see that Kalmannet infers faster than the classical methods, thanks to the highly efficient neural network computations and the fact that Kalmannet does not involve matrix inversions for each time step.

TABLE IX: Lorenz attractor with sampling mismatch.

(b) $L \equiv 20$, $\nu = 0$ [dB], $\mathbf{h}(\cdot)$ non-linear.



State-observedations sampling mismatch: We conclude our experimental study of the Lorenz attractor setup with an evaluation of KalmanNet in the presence of sampling mismatch. Here, we generate data from the Lorenz attractor SS model with an approximate continuous-time evolution process using a dense sampling rate, set to $\Delta t = 10^{-5}$. We then sub-sample the noiseless observations from the evolution process by a ratio

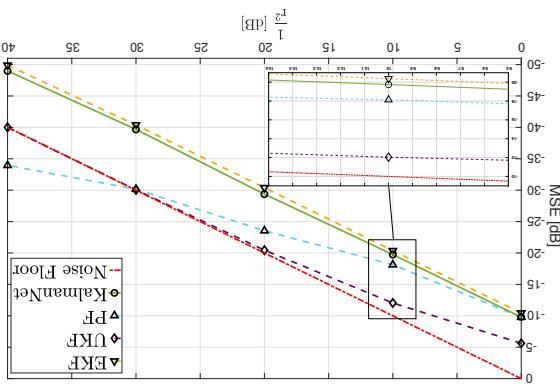
cause a severe performance degradation for the MB filters, while KalmannNet (with setting C3) is able to learn from data to overcome such mismatches and to notably outperform its MB counterparts, which are sensitive to model uncertainty. Here, we trained KalmannNet on short trajectories with $T = 100$ time steps, tested it on longer trajectories with $T = 1000$ time steps, and set $\nu = -20$ [dB]. This again demonstrates that the learning of KalmannNet is transferable.

$\theta = 0^\circ$	EKF	-10.40	-20.41	-30.50	-40.45	0.0	10	20	30
$\theta = 10^\circ$	EKF	-3.48	-6.20	-10.50	-16.50	0.0	0.35	0.37	0.34
$\theta = 20^\circ$	EKF	-18.57	-18.19	-18.00	-18.00	0.0	0.54	0.51	0.22
$\theta = 30^\circ$	EKF	-18.57	-18.19	-18.00	-18.00	0.0	0.53	0.59	0.62
$\theta = 40^\circ$	UKF	-1.18	-2.08	-2.92	-6.92	0.0	1.73	1.53	1.59
$\theta = 50^\circ$	PF	-19.87	-8.48	-0.18	15.24	0.0	8.21	3.50	0.80
$\theta = 60^\circ$	KalmanNet	-34.04	-18.17	-27.32	-27.32	0.0	0.53	0.42	0.77

TABLE VIII: MSE [dB] - Lorenz attractor with observation rotation.

$J = 2$	EKF	μ	-20.37	-30.40	-40.39	-49.89	10	20	30	40	$1/\sqrt{2} [\text{dB}]$
$J = 5$	EKF	μ	-0.25	-0.24	-0.24	-0.20	-19.47	-23.63	-33.51	-41.15	-41.74
$J = 2$	EKF	μ	-0.25	-0.25	-0.25	-0.20	-11.95	-20.45	-30.05	-39.98	-41.74
$J = 2$	UKF	μ	-0.87	-0.27	-0.27	-0.09	-0.09	-0.09	-0.09	-0.09	-41.74
$J = 2$	PF	μ	-17.95	-23.47	-30.11	-33.81	-19.71	-27.07	-35.41	-41.74	-41.74
$J = 2$	$KalmanNet$	μ	-0.18	-0.09	-0.10	-0.13	-0.09	-0.18	-0.29	-0.20	-0.11

TABLE VII: MSE [dB] - Lorenz attractor with state-evolution



In this work we presented Kalmannet, a hybrid combination of deep learning with the classic MB EKF. Our design includes the SS-model-dependent computations of the MB EKF, and the deep learning with the classic MB EKF. Our design is decoupled from the y -axis, and the linear SS model used in the assumed model, the x -axis (in cartesian coordinates)

V. CONCLUSIONS

modules for real world applications. Fig. 11 and Table X demonstrate the superiority of Kalmannet over the Kalman RNN, which is agnostic of the motion model, fails to localize. Kalmannet overcomes the errors induced by the noisy odometry observations, and provides the most accurate real-time localizations, demostating the gains of combining MB EKF-based inference with integrated DD techniques.

Fig. 11 and Table X demonstrate the superiority of Kalmannet over such scenarios. KF blindly follows the odometer trajectory and is incapable of accounting for the drift, producing a very similar or even worse estimation than the integrated RNN, which is agnostic of the motion velocity. The vanilla RNN, which is agnostic of the motion velocity, fails to localize. Kalmannet overcomes the errors induced by the noisy odometry observations, and provides the most accurate real-time localizations, demostating the gains of combining MB EKF-based inference with integrated DD techniques.

We arbitrarily use the session with date 2012-01-22 that

consists of a single trajectory. Sampling at [Hz] results in

5,850 time steps. We removed unstable readings and were

left with 5,56 time steps. The trajectory was split into three

sections: 85% for training (23 sequences of length $T = 200$),

10% for validation (2 sequences of length $T = 200$), and 5% for testing (1 sequence, $T = 277$). We compare Kalmannet with setting $\Delta\theta = 1^\circ$ to end-to-end vanilla RNN and the MB EKF, where for the

Kalmannet does not rely on knowledge of the noise covariances, \mathbf{Q} is given here for the use of the MB

EKF and for completeness.

Since Kalmannet does not rely on knowledge of the noise

covariances, \mathbf{Q} is given here for the use of the MB

EKF and \mathbf{R} are the position and velocity, respectively.

The discrete-time state-evolution with sampling interval Δt

is approximated as a linear SS model in which the evolution

matrix \mathbf{F} and noise covariance \mathbf{Q} are given by

$$\mathbf{F} = \begin{pmatrix} 1 & \Delta t \\ 0 & 1 \end{pmatrix}, \quad \mathbf{Q} = q^2 \cdot \begin{pmatrix} \frac{1}{2} & (\Delta t)^2 \\ \frac{1}{2} \cdot (\Delta t)^2 & \frac{1}{2} \end{pmatrix}. \quad (23)$$

where the acceleration is modeled as a white Gaussian noise

process w_t with variance q^2 [36].

To tackle this problem we model the Segway kinematics (in

TABLE X: Numerical MSE [dB] for the NCLT experiment.

Baseline	EKF	Kalmannet	Vanilla RNN
25.47	25.385	22.2	40.21

for Kalmann filtering is given by

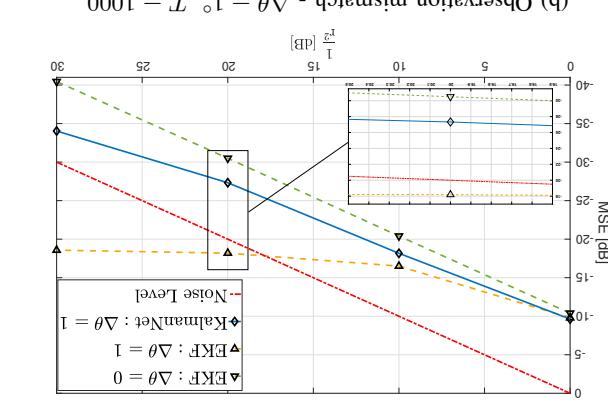


Fig. 9: Lorenz attractor, partial information.

are decoupled from the y -axis, and the linear SS model used in the assumed model, the x -axis (in cartesian coordinates)

estimated positions typically start drifting away at some point. Such settings where one does not have access to direct measurements for positioning are very challenging yet practical and typical for many applications where positioning noisy odometry readings are not available indoors, and one must rely on noisy odometry readings for self-localization. Odometry-based techniques are not available indoors, and one must rely on noisy velocity readings. In this case the observations obey a noisy linear model:

$$y \in \mathbb{R}, \quad H = (0, 1). \quad (24)$$

The goal is to track the underlying state vector in both axes solely using odometry data; i.e., the observations are given by noisy velocity readings. In this case the observations obey a noisy linear model:

$$y \in \mathbb{R}, \quad H = (0, 1). \quad (24)$$

Since Kalmannet does not rely on knowledge of the noise covariances, \mathbf{Q} is given here for the use of the MB

EKF and \mathbf{R} are the position and velocity, respectively.

The discrete-time state-evolution with sampling interval Δt

is approximated as a linear SS model in which the evolution

matrix \mathbf{F} and noise covariance \mathbf{Q} are given by

$$\mathbf{F} = (p, v)_T \in \mathbb{R}^2, \quad \frac{\partial}{\partial t} \mathbf{x}_t = \begin{pmatrix} 0 & 0 \\ 0 & 1 \end{pmatrix} \cdot \mathbf{x}_t + \begin{pmatrix} 0 \\ 0 \end{pmatrix}. \quad (22)$$

where the acceleration is modeled as a white Gaussian noise

process w_t with variance q^2 [36].

To tackle this problem we model the Segway kinematics (in

readings (e.g., GPS and odometer) and the ground truth loca-

tions of a moving Segway robot. Given these noisy readings,

the goal of the tracking algorithm is to localize the Segway

from the raw measurements at any given time.

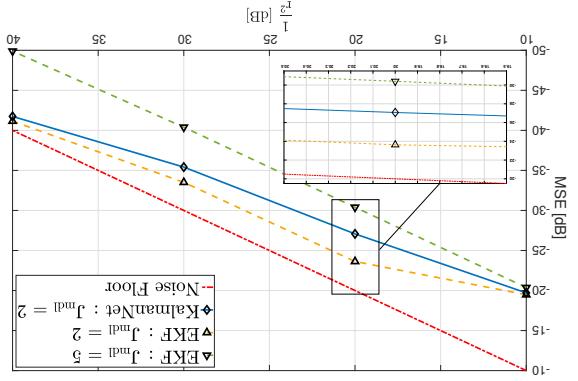
readings (e.g., GPS and odometer) and the ground truth loca-

tions of a moving Segway robot. Given these noisy readings,

the goal of the tracking algorithm is to localize the Segway

from the raw measurements at any given time.

(a) State-evolution mismatch, identity h , $T = 2000$.



- [11] N. J. Gordon, D. J. Salmond, and A. F. Smith, "Novel approach to noninertial Non-Gaussian Bayesian state estimation," in *IEEE proceedings F (radar and signal processing)*, vol. 140, no. 2, pp. 107-113, 1998.
- [12] P. Del Moral, "Nonlinear filtering: Interacting particle resolution," *Comptes Rendus de l'Academie des Sciences - Series I-Mathematics*, vol. 325, no. 6, pp. 653-658, 1997.
- [13] J. S. Liu and R. Chen, "Sequential Monte Carlo methods for dynamic systems," *Journal of the American Statistical Association*, vol. 93, no. 443, pp. 102-104, 1998.
- [14] F. Auger, M. Hilario, J. M. Guerrero, E. Monmasson, T. Orlowska-Kowalska, and S. Krasura, "Industrial application of the Kalman filter: A review," *IEEE Trans. Ind. Electronics*, vol. 60, no. 12, pp. 5458-5471, 2013.
- [15] M. Zorzi, "Robust Kalman filtering under model perturbations," *IEEE Trans. Autom. Control*, vol. 62, no. 6, pp. 2902-2907, 2016.
- [16] ——, "On the robustness of the Bayes and Wiener estimators under model uncertainty," *Automatica*, vol. 133, pp. 133-140, 2017.
- [17] A. Longtin, M. Perellin, S. Gorodetsky, S. Yi, H. Liu, and M. Zorzi, "Learning the tuned liquid damper dynamics by means of a robust EKF," *arXiv preprint arXiv:2103.03220*, 2021.
- [18] Y. LeCun, Y. Bengio, C. Burcin, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, p. 436, 2015.
- [19] Y. Bengio, "Learning deep architectures for AI," *Foundations and Trends® in Machine Learning*, vol. 2, no. 1, pp. 1-127, 2009.
- [20] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735-1780, 1997.
- [21] C. Chung, C. Gliech, C. Schmidhuber, "Empirical evaluation of gated recurrent neural networks on sequence modeling," *arXiv preprint arXiv:1412.3555*, 2014.
- [22] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," *arXiv preprint arXiv:1706.03762*, 2017.
- [23] M. Zahedi, A. Almude, and A. J. Smola, "Latent LSTM allocation: *arXiv:1706.03762*, 2017.
- [24] N. Shlezinger, N. Farstad, Y. C. Eldar, "DeepSIC: Deep soft interference cancellation for multiuser MIMO detection," *IEEE Trans. Wireless Communications*, vol. 19, no. 5, pp. 3319-3331, 2020.
- [25] N. Shlezinger, R. Fu, and Y. C. Eldar, "DeepSIC: Deep soft interference cancellation for multiuser MIMO detection," *IEEE Trans. Wireless Communications*, vol. 20, no. 2, pp. 1349-1362, 2021.
- [26] N. Shlezinger, N. Farstad, Y. C. Eldar, and A. J. Goldsmith, "Learned feature graphs for inference from noisy long-term visual sequences," *IEEE Trans. Signal Processing*, early access, 2022.
- [27] N. Shlezinger, J. Shabtai, and D. Sontag, "Deep Kalman filters," *arXiv preprint arXiv:1511.05121*, 2015.
- [28] N. Carreira-Perpinan, U. Paquet, and D. Blei, "Variational deep learning," *arXiv preprint arXiv:1605.06432*, 2016.
- [29] R. G. Krishnamoorthy, S. D. Kamroni, U. Paquet, and D. Blei, "A disentangled variational autoencoder and variational Bayes filters: Unsupervised learning of state space models from raw data," *arXiv preprint arXiv:1605.06432*, 2016.
- [30] M. Karthik, M. Soelech, J. Baye, and P. Van der Smagt, "Deep variational Bayes filters: Unsupervised learning of state space models from raw data," *arXiv preprint arXiv:2012.08405*, 2020.
- [31] M. Fraccaro, S. D. Kamroni, U. Paquet, and D. Blei, "A disentangled variational autoencoder and variational Bayes filters in *Neural Information Processing Systems*," *MLR*, 2018, pp. 968-977.
- [32] C. Neesseth, S. Lindemann, R. Rangamathan, and D. Blei, "Variational inference in *Neural Information Processing Systems*," in *Inferential Conference and Statistics*, pp. 968-977.

- [33] J. D. Meléndez, S. F. Schmidt, and L. A. McGee, "Optimal filtering and linear prediction applied to a midcourse trajectory estimation," *Stanford Research Institute Tech. Rep.*, 1967.
- [34] R. E. Larson, R. M. Dressler, and R. S. Ratner, "Application of the Extended Kalman filter to ballistic trajectory estimation," *MIT Lincoln Lab. Tech. Rep.*, 1967.
- [35] M. Grinber, "An approach to target tracking," *MIT Lexington Lincoln MA*, 1949, vol. 8.
- [36] N. Wiener, *Extrapolation, interpolation, and smoothing of stationary time series: With engineering applications*. MIT Press Cambridge, 1949.
- [37] R. E. Kalman, "New methods in Wiener filtering theory," 1963.
- [38] J. Durbin and S. J. Koopman, *The time series analysis by state space methods*. Oxford University Press, 2012.
- [39] R. E. Kalman, "A new approach to linear filtering and prediction problems," *Journal of Basic Engineering*, vol. 82, no. 1, pp. 35-45, 1960.
- [40] R. E. Kalman and R. S. Bucy, "New results in linear filtering and prediction theory," *Journal of Basic Engineering*, vol. 83, no. 1, pp. 96-107, 1961.
- [41] G. Breach, N. Shlezinger, R. J. G. van Slooten, and Y. C. Eldar, "KalmanNet: Data-driven Kalman filtering," in *Proc. IEEE ICASSP 2021*, pp. 3905-3909.
- [42] J. Durbin and S. J. Koopman, *The time series analysis by state space methods*. Oxford University Press, 2012.
- [43] R. E. Kalman, "A new approach to linear filtering and prediction problems," *Journal of Basic Engineering*, vol. 82, no. 1, pp. 35-45, 1960.
- [44] R. E. Kalman and R. S. Bucy, "New results in linear filtering and prediction theory," *Journal of Basic Engineering*, vol. 83, no. 1, pp. 96-107, 1961.
- [45] J. Durbin and S. J. Koopman, *The time series analysis by state space methods*. Oxford University Press, 2012.
- [46] N. Wiener, *Extrapolation, interpolation, and smoothing of stationary time series: With engineering applications*. MIT Press Cambridge, 1949.
- [47] M. Grinber, "An approach to target tracking," *MIT Lexington Lincoln MA*, 1949, vol. 8.
- [48] J. Durbin and S. J. Koopman, *The time series analysis by state space methods*. Oxford University Press, 2012.
- [49] R. E. Kalman, "A new approach to linear filtering and prediction problems," *Journal of Basic Engineering*, vol. 82, no. 1, pp. 35-45, 1960.
- [50] J. Durbin and S. J. Koopman, *The time series analysis by state space methods*. Oxford University Press, 2012.
- [51] R. E. Kalman, "New methods in Wiener filtering theory," 1963.

REFERENCES

We would like to thank Prof. Hans-Andrea Loeliger for his helpful comments and discussions, and Jonas E. Meier for his assistance with the numerical study.

ACKNOWLEDGMENTS

Fig. 11: NC LT data set: ground truth vs. integrated velocity features encapsulating the information needed for its operation. Our numerical study shows that doing so enables KalmanNet to carry out real-time state estimation in the same manner as MB Kalman filter, while learning to overcome model mismatches and non-linearities. KalmanNet uses a relatively small compact RNN that can be trained with a relatively small data set and infers a reduced complexity, making it applicable for high dimensional SS models and computationally limited devices.

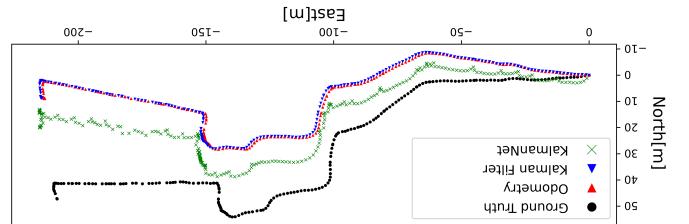
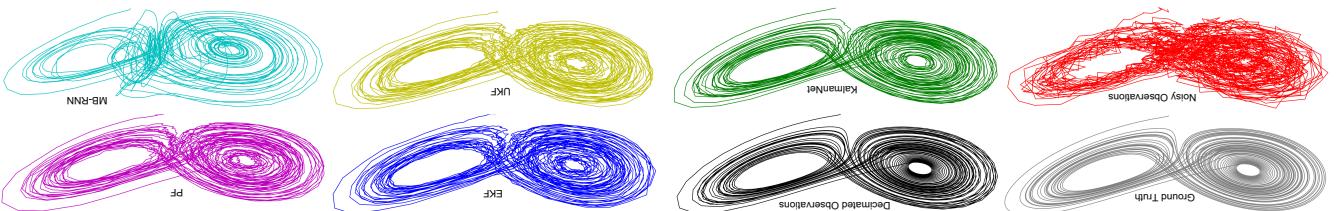


Fig. 10: Lorenz attractor with sampling mismatch (decimation), $T = 3000$.



- [33] E. Archer, I. M. Park, L. Buesing, J. Cuningham, and L. Paniaski, “Black box variational inference for state space models”, *arXiv preprint arXiv:1511.07367*, 2015.
- [34] R. Krishnamoorthy, U. Shalil, and D. Soniatte, “Structured inference networks for non-linear state space models”, in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 31, no. 1, 2017.
- [35] G. Satorras, Z. Akata, and M. Weillige, “Combining generative and discriminative models for hybrid inference”, in *Advances in Neural Information Processing Systems*, 2016, pp. 1382–1382.
- [36] Y. Bar-Schalom, X. R. Li, and T. Kirubarajan, “Estimation with applications to tracking and navigation: Theory, algorithms and software”, John Wiley & Sons, 2004.
- [37] K.-V. Yuen and S.-C. Kuok, “Online updating and uncertainty quantification using noisy measurements for output-only measurement”, *Mechanical Systems and Signal Processing*, vol. 66, pp. 62–77, 2016.
- [38] K.-V. Yuen and S.-C. Kuok, “Stable robust Extended Kalman filter using noisy measurements for output-only measurement”, *Journal of Aerospace Engineering*, vol. 30, no. 2, p. B416010, 2017.
- [39] I. Arasaratnam, S. Haykin, and R. J. Elliott, “Discrete-time nonlinear filtering algorithms using Gauss-Hermite quadrature”, *Proc. IEEE*, vol. 95, no. 5, pp. 953–977, 2007.
- [40] I. Arasaratnam and S. Haykin, “Cubature Kalman filters”, *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 75, no. 3, pp. 397–426, 2013.
- [41] M. Arulampalam, S. Maskell, N. Gordon, and T. Clapp, “A tutorial on particle filters for nonlinear state-space modeling”, *IEEE Transactions on Signal Processing*, vol. 51, no. 3, pp. 1243–1269, 2003.
- [42] N. Chopin, P. E. Jacob, and O. Papaspiliopoulos, “SMC2: An efficient algorithm for particle filters for multivariate Gaussian-Non-Gaussian tracking”, *IEEE Transactions on Signal Processing*, vol. 50, no. 2, pp. 174–188, 2002.
- [43] L. Martino, V. Elvira, and G. Camps-Valls, “Distributed particle filters for sequential analysis of state space models”, *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 75, no. 3, pp. 553–557.
- [44] C. Andrieu, A. Doucet, and R. Holenstein, “Particle Markov chain Monte Carlo methods”, *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 72, no. 3, pp. 269–342, 2010.
- [45] J. Effring, E. Tora, and R. van de Molengraft, “Particle filters: A hands-on tutorial”, *Sensors*, vol. 21, no. 2, p. 438, 2021.
- [46] A. Hurniwy and D. S. Stoilescu, “An approach to time series smoothing and forecasting using the EM algorithm”, *Journal of Time Series Analysis*, vol. 3, no. 4, pp. 233–264, 1982.
- [47] Z. Ghahramani and G. E. Hinton, “Parameter estimation for linear dynamical systems”, 1996.
- [48] J. Daubert, A. Eckford, S. Kot, and H.-A. Loeliger, “Expectation maximization as message passing-part I: Principles and Gaussian max-products”, *IEEE Transactions on Signal Processing*, vol. 50, no. 1, pp. 172–185, 2002.
- [49] L. Martino, J. Reed, V. Elvira, and F. Louzada, “Cooperative parallel particle filters for online model selection and applications to urban mobility”, *Digital Signal Processing*, vol. 69, pp. 1526–1531.
- [50] P. Abbeel, A. Coates, M. Moninger, A. Ng, and S. Thrun, “Discriminative training of Kalman filters”, in *Robotics: Science and Systems*, vol. 2, 2005, p. 1.
- [51] L. Xu and R. Niu, “ERFNet: Learning noise statistics from data”, in *Proc. IEEE ICASSP*, 2021, pp. 4560–4564.
- [52] S. T. Barrau and S. Boyd, “Robust Kalman smoother for learning discrete-time systems”, *IEEE Trans. Autom. Control*, vol. 39, no. 6, pp. 1310–1314, 1994.
- [53] J. Xie, Y. C. Soh, and C. E. De Souza, “Robust Kalman filtering for uncertain and smooth systems”, *IEEE Trans. Autom. Control*, vol. 39, no. 6, pp. 1310–1314, 1994.
- [54] C. M. Carvalho, M. S. Jofmanns, H. F. Lopes, and N. G. Polson, “Particule 2010. Monte Carlo Control Conference, vol. 25, no. 1, pp. 88–106,
- [55] I. Uruega, M. F. Bugarito, and P. M. Djuric, “Sequential Monte Carlo Workshop (SSP), 2016, pp. 1–5.
- [56] L. Zhou, Z. Luo, T. Shen, J. Zhang, M. Zhen, Y. Yao, T. Fang, and L. Guan, “KFNet: Learning emporial camera relocalization using and deep generative models”, *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 4919–4928.
- [57] D. J. Rezende, S. Mohamed, and D. Wierstra, “Stochastic backpropagation and approximate inference in deep generative models”, *arXiv preprint arXiv:1312.6114*, 2013.
- [58] L. Zhou, Z. Luo, T. Shen, J. Zhang, M. Zhen, Y. Yao, T. Fang, and L. Guan, “KFNet: Learning emporial camera relocalization using and deep generative models”, *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 4919–4928.
- [59] D. M. Blei, A. Kucukelbir, and J. D. McAuliffe, “Variational inference: A review for statisticians”, *Journal of the American Statistical Association*, vol. 112, no. 518, pp. 859–877, 2017.
- [60] T. Haranosa, A. Ajay, S. Levine, and P. Abbeel, “Backprop kf: Learning dynamics via tracking state space models”, *Journal of the IEEE International Conference on Neural Information Processing Systems*, 2018, pp. 7785–7794.
- [61] B. Larner-Goldschtein, R. Tlmon, and S. Gamal, “A hybrid approach for short-term memory Kalman filters: Recurrent neural estimators for pose regularization”, *Proc. IEEE International Conference on Machine Learning*, 2017.
- [62] H. Coskun, F. Achille, R. Dipietro, N. Navab, and F. Tombari, “Long-term memory in Neural Information Processing Systems”, 2018, pp. 7785–7794.
- [63] S. S. Ramaprabhan, M. W. Seeger, J. M. Gasparian, L. Seelli, Y. Wang, and G. Neumann, “Recurrent Kalman networks: Factorized inferred inference in high-dimensional deep feature spaces”, in *International Conference on Machine Learning*, 2019, pp. 544–552.
- [64] P. Beekker, H. Pandya, G. Geethardhi, C. Zhao, C. J. Taylor, and T. Zhang, “State space LSTM models with particle MCMC inference”, *ICASSP*, 2022.
- [65] X. Ni, G. Revach, N. Shlezinger, J. E. Meir, R. J. van Sloun, and Y. Eldar, “Unertainty in data-driven Kalman filtering for parallel known state-space models”, in *Proc. IEEE ICASSP*, 2022.
- [66] T. Salimans, D. Kingma, and M. Welling, “Markov chain Monte Carlo and variational inference”, *Proc. NIPS*, 2017.
- [67] X. Ni, G. Revach, N. Shlezinger, J. E. Meir, R. J. van Sloun, and Y. Eldar, “Data-driven Kalman-based velocity estimation for autonomous racing”, in *Proc. IEEE ICAS*, 2021.
- [68] R. Dey and F. M. Saleem, “Gated recurrent unit (GRU) R-TSNET: Deep learning aided Kalman filters”, in *Proc. IEEE ICASSP*, 2017.
- [69] P. J. Werbos, “Backpropagation through time: What it does and how to do it”, *Proc. IEEE*, vol. 78, no. 10, pp. 1539–1560, 1990.
- [70] I. Stookeve, “Training recurrent neural networks”, University of Toronto, Canada, 2013.
- [71] I. Kefin, G. Revach, N. Shlezinger, J. E. Meir, R. J. van Sloun, and Y. Eldar, “Unertainty in data-driven Kalman filtering for parallel learning”, *Proc. ICASSP*, 2013.
- [72] A. Lopez Escoriza, G. Revach, N. Shlezinger, J. E. Meir, R. J. van Sloun, and Y. Eldar, “Unertainty in data-driven Kalman filtering for parallel learning”, *Proc. ICASSP*, 2022.
- [73] H. E. Rauch, F. Tung, and C. T. Striebel, “Maximum likelihood estimation of linear dynamic systems”, *AIAA Journal*, vol. 3, no. 8, pp. 1445–1450, 1965.
- [74] P. Kriegsmann and J. Ba, “Adam: A method for stochastic optimization”, *arXiv preprint arXiv:1412.6980*, 2014.
- [75] Labbe, Roger, Filterpy - Kalman and Bayesian Filters in Python, 2020.
- [76] Jecker Nordahl, *PyParticleFits - Particle based methods in Python*, 2015.
- [77] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [78] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [79] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [80] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [81] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [82] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [83] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [84] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [85] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [86] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [87] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [88] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [89] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [90] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [91] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [92] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [93] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [94] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [95] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [96] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [97] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [98] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [99] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [100] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [101] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [102] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [103] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [104] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [105] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [106] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [107] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [108] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [109] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [110] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [111] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [112] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [113] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [114] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [115] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [116] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [117] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [118] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [119] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [120] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [121] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [122] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [123] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [124] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [125] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [126] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [127] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [128] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [129] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [130] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [131] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [132] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [133] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [134] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [135] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [136] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [137] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [138] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [139] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [140] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [141] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [142] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [143] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [144] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [145] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [146] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [147] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [148] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [149] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [150] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [151] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [152] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [153] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [154] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [155] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [156] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [157] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [158] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [159] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [160] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [161] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [162] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [163] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [164] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [165] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [166] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [167] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [168] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [169] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [170] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [171] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [172] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [173] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [174] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [175] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [176] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [177] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [178] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [179] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [180] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [181] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [182] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [183] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [184] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [185] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [186] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [187] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [188] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [189] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [190] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [191] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [192] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [193] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [194] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [195] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [196] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [197] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [198] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [199] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [200] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [201] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [202] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [203] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [204] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [205] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [206] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [207] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [208] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [209] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [210] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [211] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [212] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [213] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [214] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [215] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [216] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [217] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [218] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [219] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [220] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [221] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [222] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [223] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [224] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [225] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [226] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [227] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [228] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [229] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [230] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [231] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [232] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [233] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [234] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [235] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [236] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [237] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [238] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [239] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [240] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [241] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [242] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [243] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [244] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [245] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [246] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [247] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [248] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [249] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [250] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [251] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [252] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [253] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [254] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [255] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [256] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [257] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [258] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [259] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [260] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [261] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [262] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [263] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [264] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [265] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [266] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [267] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [268] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [269] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [270] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [271] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [272] [Online]. Available: <https://pypi.readthedocs.io/en/latest/>
- [273] [Online]. Available: [https://pypi.readthedocs.io/en/latest/</a](https://pypi.readthedocs.io/en/latest/)